

# AI Developer Task: RAG-based Chatbot (Company Policies)

## Objective

Build a **Retrieval-Augmented Generation (RAG)** chatbot that can answer questions related to **company policies**.

This task is designed to test your ability to work with **LLMs, embeddings, and vector databases** — key skills for AI development.

---

## Task Description

Create a chatbot that:

1. Ingests and stores a few text or PDF files containing **sample company policies** (you can use public HR or IT policy templates if needed).
  2. Uses **retrieval + generation** to answer user queries — e.g.:
    - “What is the company’s leave policy?”
    - “How can employees request remote work?”
  3. Retrieves the most relevant document chunks using **embeddings + vector search**.
  4. Generates a natural, concise response **based only on retrieved context** (no hallucination).
  5. Shows the **retrieved source text or document name** as citation.
- 

## Tech Stack (Recommended)

You can choose your preferred tools, but the recommended options are:

- **LLM:** OpenAI GPT, Llama 3, Gemini, Groq, etc.
- **Embeddings:** OpenAI, HuggingFace, SentenceTransformers, Cohere, Gemini
- **Vector Store:** FAISS, Chroma, or Pinecone
- **Framework:** Any suitable backend framework
- **Interface:** Streamlit, Gradio, or Swagger API endpoints

**Note:** You may also use your **own preferred tools or setup**.

---

## Bonus Points

- Add conversation history / memory
- 

## Submission Format

Please submit:

- A **GitHub repository** containing:
    - Your code
    - A **README.md** explaining setup steps and approach
  - Optional: A **demo video** (1–2 min)
- 

## Time Limit

You have **48 hours** to complete and submit your solution after receiving the task.

---

## Note

If you face any **challenges** or need **clarification** during the task, feel free to **communicate via email** at [simnsnto@gmail.com](mailto:simnsnto@gmail.com) or [s.shahriar32322323@gmail.com](mailto:s.shahriar32322323@gmail.com) for assistance.